ML-First ABM Anomaly Detection System - Complete Download Package

What You Get

A complete, production-ready ML-first anomaly detection system for ABM Electronic Journals with:

- Pure ML Detection: BERT embeddings + unsupervised learning (no regex parsing first)
- Expert Labeling Interface: Web UI for domain experts to label anomalies
- Supervised Learning: Trains on expert-labeled data for improved accuracy
- Real-time Dashboard: Live monitoring with beautiful visualizations
- **Docker Deployment**: Everything containerized and ready to run

Quick Deployment (10 Minutes)

Step 1: Download and Setup

Save and run the setup script:

```
bash
```

```
# Save the setup script from "Complete ML-First ABM Anomaly Detection Setup Script" chmod +x setup_ml_first_abm_system.sh
./setup_ml_first_abm_system.sh
```

Step 2: Integrate ML Components

```
bash
```

```
cd abm-anomaly-ml-first
chmod +x integrate_ml_components.sh
./integrate_ml_components.sh
```

This adds all the ML code, API endpoints, and UI components automatically.

Step 3: Build and Deploy

bash

Build all Docker images (this takes ~5-10 minutes first time) make build

Start all services

make up

Verify everything is running make components docker-compose ps

Step 4: Generate Test Data

bash

Create test ABM logs with various anomalies
python3 generate_test_abm_logs.py

Copy to input directory for processing
cp test_abm_logs.txt data/input/

Access Your System

Within 2-3 minutes, you can access:

Service	URL	Purpose
Dashboard	http://localhost:3000	Main monitoring interface
Expert Labeling	http://localhost:3000	Click "Expert Review" tab
API Docs	http://localhost:8000/docs	Interactive API documentation
Jupyter	http://localhost:8888	Data analysis (token: ml_jupyter_token_123)
Grafana	http://localhost:3001	System monitoring (login: admin) (ml_admin)

■ How to Use

1. Process EJ Logs

Option A: Upload via Dashboard

- 1. Click "Upload EJournal" button
- 2. Select your log file
- 3. Processing starts automatically

Option B: Drop in folder

cp your_ej_logs.txt data/input/

Files are processed automatically every minute

2. Review Detected Anomalies

- 1. Go to Dashboard → Anomalies tab
- 2. See all ML-detected anomalies with scores
- 3. View patterns found without regex

3. Expert Labeling Process

- 1. Navigate to "Expert Review" tab
- 2. Review each anomaly:
 - See raw log text
 - View detected patterns
 - Check anomaly score
- 3. Label anomalies:
 - Select from predefined labels
 - Add custom labels
 - Mark false positives as "Not Anomaly"
- 4. Click "Save All Labels"

4. Train Supervised Model

Once you have labeled at least 10 anomalies:

- 1. Click "Train Supervised Model" button
- 2. Wait for training to complete (~1-2 minutes)
- 3. Future detections will use both unsupervised + supervised models

Q Understanding the ML-First Approach

Traditional (Regex-First) Flow:

Raw Logs → Regex Parsing → Structured Data → ML Models → Anomalies



X Rigid patterns miss unknown anomalies

Our ML-First Flow:



Understands context, finds new patterns

What Makes This Special

- 1. No Regex Required: Works on raw, unstructured logs
- 2. **Discovers Unknown Patterns**: Not limited to predefined rules
- 3. Continuous Learning: Gets smarter with expert feedback
- 4. **Production Ready**: Handles thousands of transactions per minute
- 5. Explainable: Shows why anomalies were detected

Solution Example Anomalies It Detects

From your requirements document:

- Unable to Dispense: (UNABLE TO DISPENSE) after normal transaction flow
- Supervisor Mode Issues: (SUPERVISOR MODE ENTRY) after transaction end
- Power Reset Problems: POWER-UP/RESET immediately after transactions
- Cash Retraction Errors: Complex (CASHIN RETRACT STARTED) patterns
- Note Handling Delays: Long gaps between (NOTES PRESENTED) and (NOTES TAKEN)

X Configuration

env

Environment Variables (.env)

```
# Customize detection sensitivity

ANOMALY_THRESHOLD=0.7 # 0.0-1.0, lower = more sensitive

# Change BERT model (for different languages)

BERT_MODEL=bert-base-uncased # or bert-base-multilingual-cased

# Database credentials (change for production!)
```

POSTGRES_PASSWORD=your_secure_password REDIS_PASSWORD=your_redis_password

Scaling for Production

```
# docker-compose.override.yml
services:
anomaly-detector:
deploy:
replicas: 3 # Run multiple instances
resources:
limits:
```

memory: 8G # Increase for larger models

Monitoring Performance

Check ML Model Status

bash

View Anomaly Statistics

bash

curl http://localhost:8000/api/v1/dashboard/stats | jq .

Export Labeled Data

bash

```
curl http://localhost:8000/api/v1/expert/export-labels?format=csv \
    -o labeled_anomalies.csv
```

Troubleshooting

Issue: Services won't start

```
bash

# Check logs

docker-compose logs -f anomaly-detector

# Verify ports are free

Isof -i :3000,8000,5432,6379
```

Issue: BERT model download fails

```
# Manually download in container

docker exec -it abm-ml-anomaly-detector python -c "
from transformers import BertModel, BertTokenizer
BertTokenizer.from_pretrained('bert-base-uncased')
BertModel.from_pretrained('bert-base-uncased')
"
```

Issue: No anomalies detected

1. Check if models are trained:

bash

```
docker exec -it abm-ml-anomaly-detector Is -la /app/models/
```

- 2. Lower the anomaly threshold in (.env)
- 3. Ensure test data has actual anomalies

Success Checklist

All containers show "Up" status in (docker-compose ps)
 Dashboard loads at http://localhost:3000
 Test file processes and shows anomalies
 Expert labeling interface works
 Can train supervised model after labeling
 Real-time alerts appear in dashboard

What's Included

```
abm-anomaly-ml-first/
---- services/
 anomaly-detector/ # ML detection engine with BERT
            # FastAPI with expert endpoints
     — aрі/
     — dashboard/ # React UI with labeling interface
   —— jupyter/ # Analysis notebooks
 —— data/
  input/ # Drop EJ logs here
                 # Trained ML models
 ---- models/
   —— sessions/ # Processed session storage
    — docker-compose.yml # Complete infrastructure
 —— Makefile # Convenience commands
    — README.md
                     # Documentation
```

Next Steps

- 1. Process Real Data: Upload your actual ABM EJ logs
- 2. Label Anomalies: Build your knowledge base
- 3. **Train Models**: Improve accuracy for your specific patterns
- 4. Set Alerts: Configure Grafana for notifications
- 5. **Scale Up**: Deploy to production with Kubernetes

Pro Tips

- 1. Start with unlabeled review: Let ML find patterns first
- 2. Label consistently: Create labeling guidelines for your team
- 3. Retrain regularly: Schedule weekly model updates
- 4. Monitor drift: Watch if anomaly patterns change over time
- 5. Export insights: Use Jupyter for deep analysis

Support

- Logs: (make logs) or (docker-compose logs [service])
- Health Check: (curl http://localhost:8000/api/v1/health)
- Documentation: Check (docs/ML_FIRST_ARCHITECTURE.md)

You now have a complete ML-first anomaly detection system that:

- V Processes raw ABM logs without regex
- V Uses BERT to understand log semantics
- V Discovers unknown anomaly patterns
- V Learns from expert feedback
- Continuously improves accuracy

Total setup time: ~10 minutes Value: Enterprise-grade ML anomaly detection

Happy anomaly hunting with ML! of i